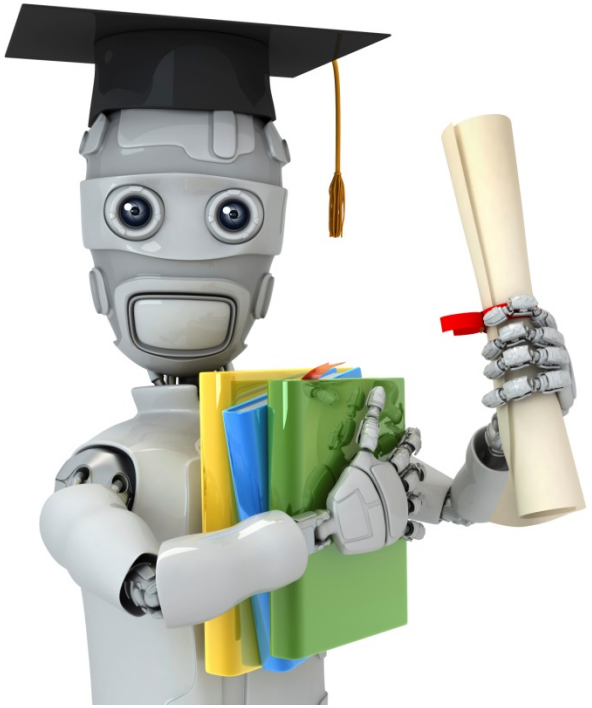

Machine Learning

Application example: Photo OCR

Problem description and pipeline

The Photo OCR problem

how to get the computer/camera to read the text/pictures better

Photo OCR pipeline

1. Text detection

2. Character segmentation

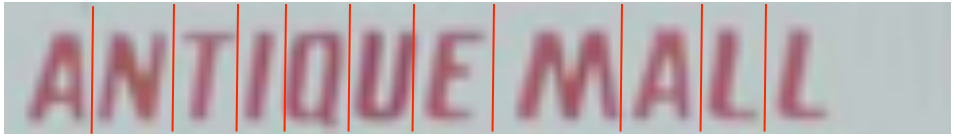

3. Character classification

Cleaning ⇒ Cleaning

Photo OCR pipeline

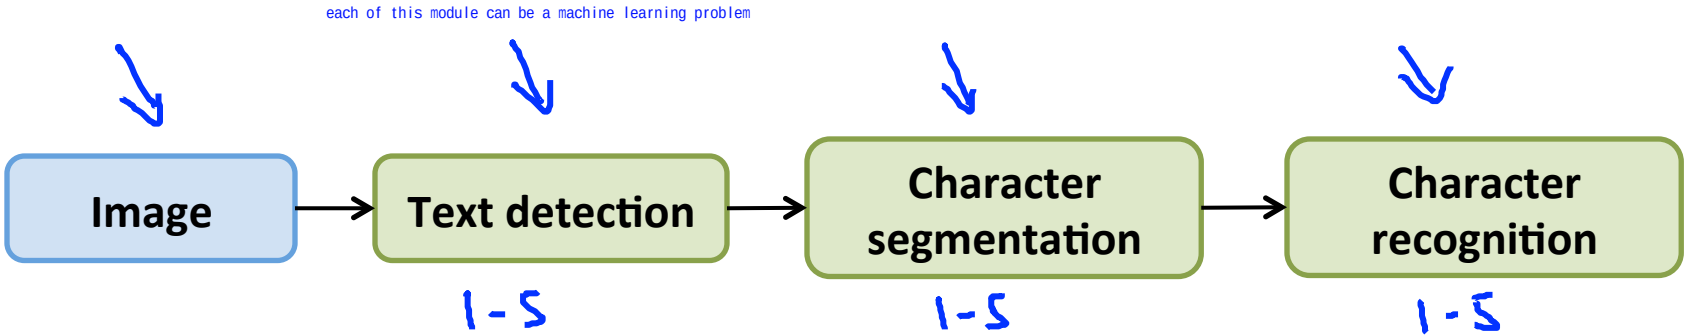

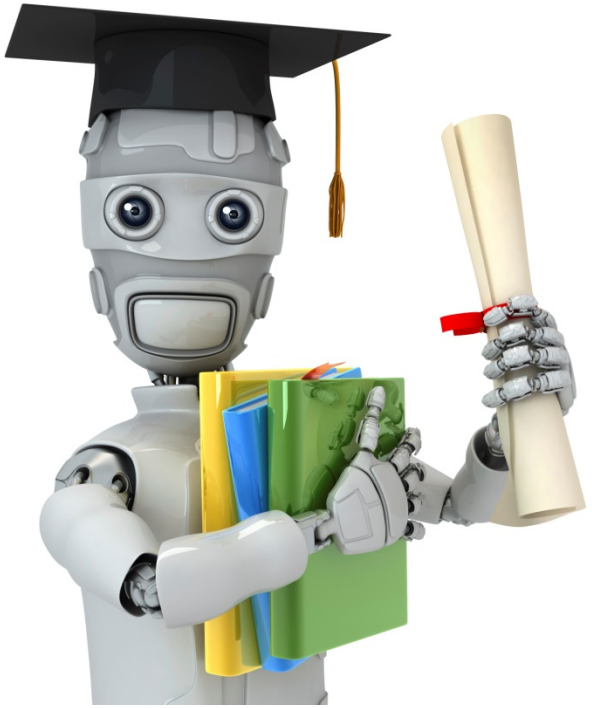

Machine Learning

Application example: Photo OCR

Sliding windows

Text detection

A simpler case

Pedestrian detection

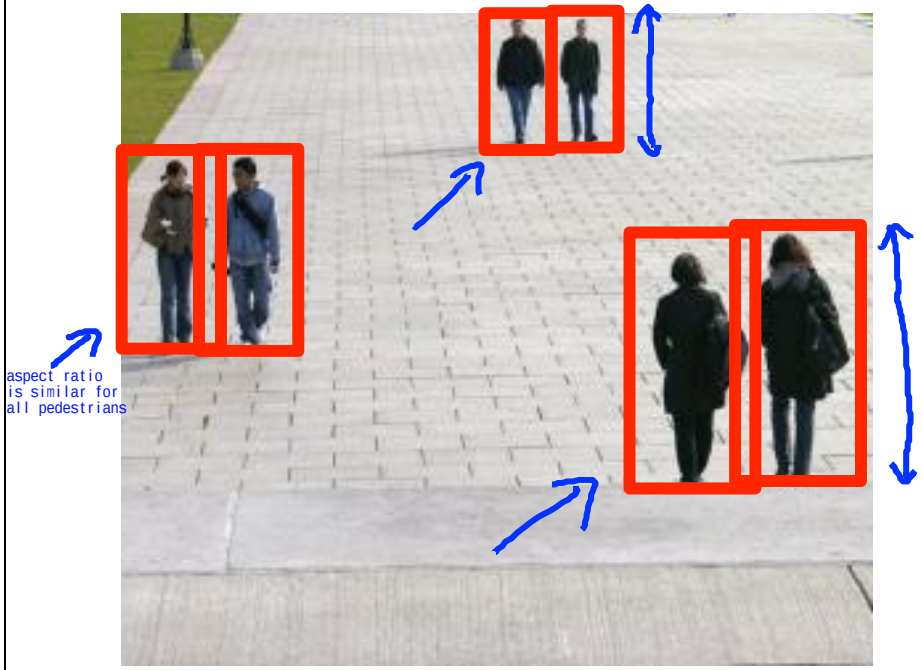

Supervised learning for pedestrian detection

$x =$ pixels in 82x36 image patches

Positive examples $(y = 1)$

Negative examples $(y = 0)$

We can use supervised learning algorithm

Sliding window detection

Sliding window detection

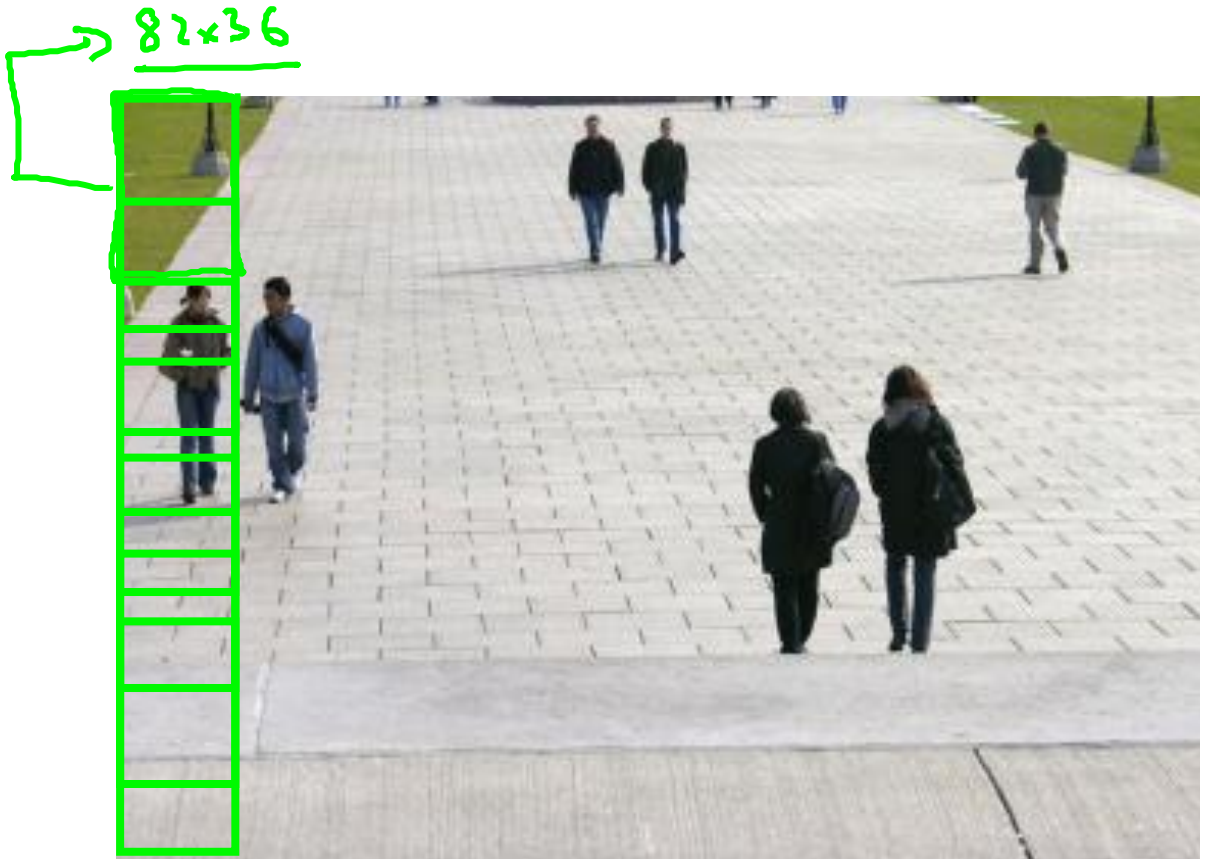

Sliding window detection

Sliding window detection

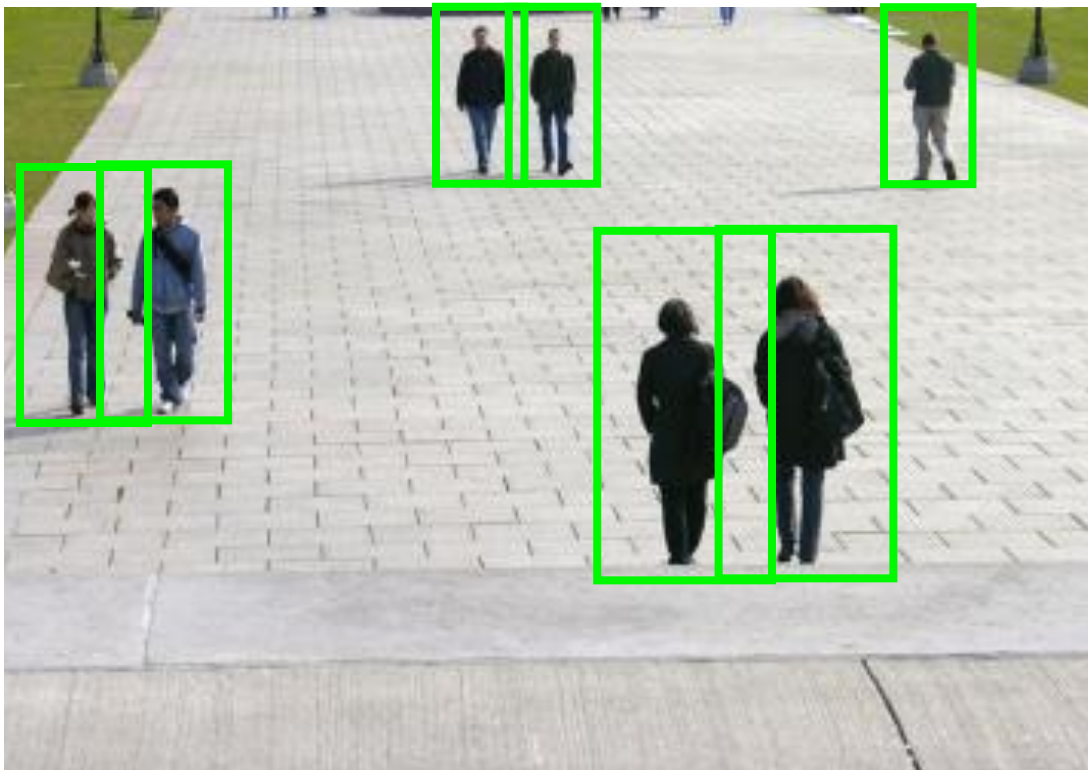

Text detection

Now instead of detecting pedestrians, we try to detect text

Text detection

Positive examples $(y = 1)$

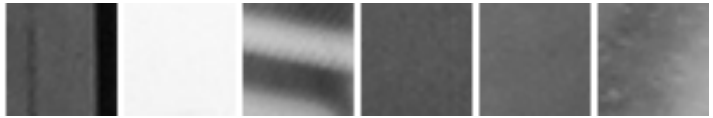

Negative examples $(y = 0)$

Text detection

"expansion"

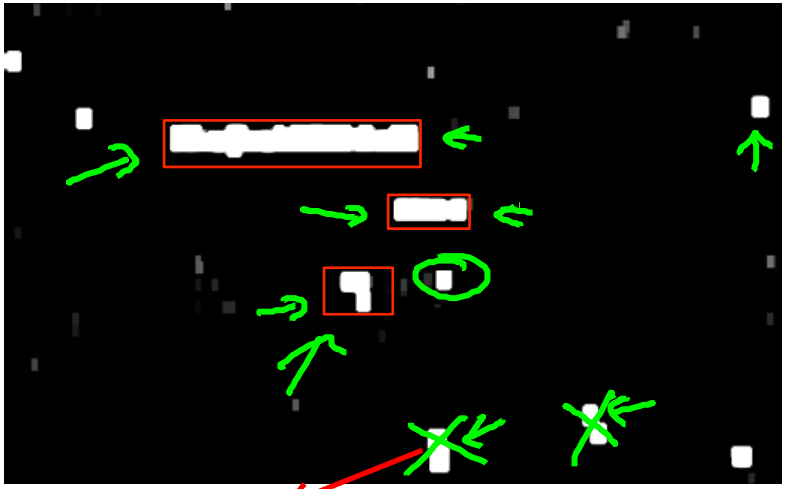

[David Wu]

Andrew Ng

1D Sliding window for character segmentation

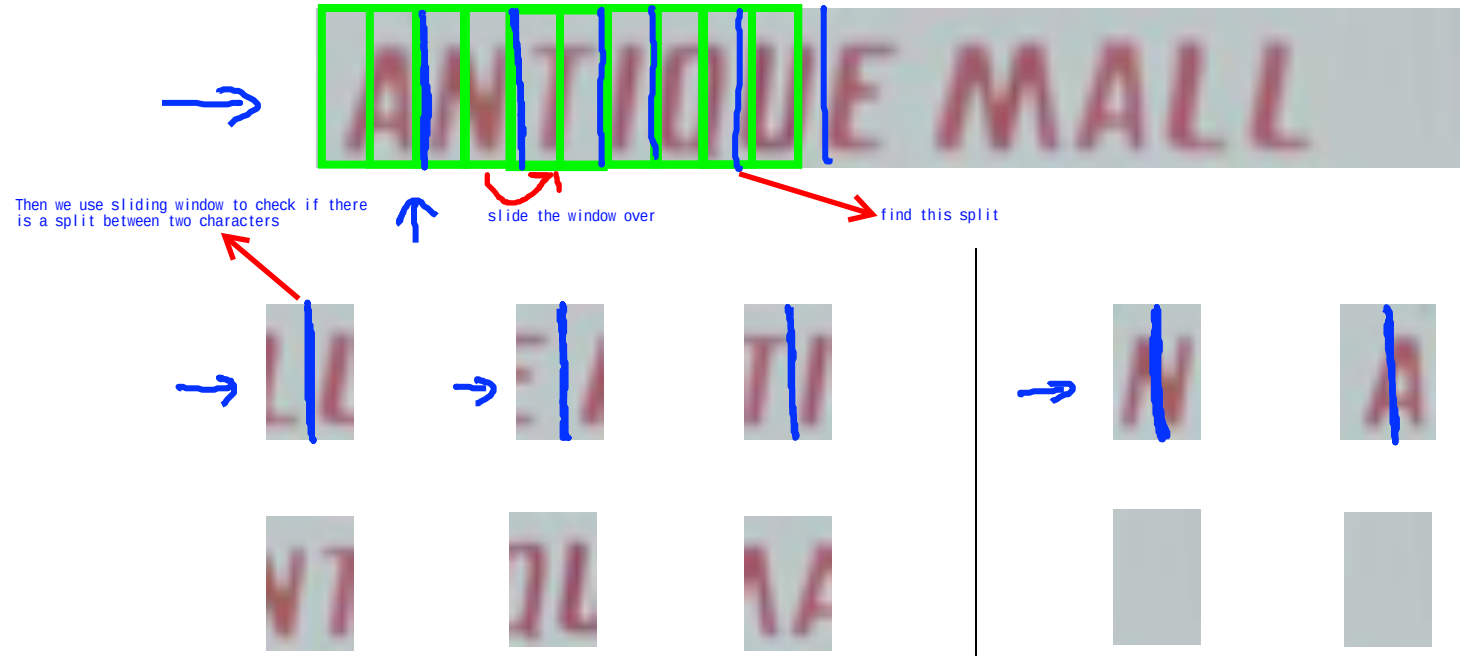

Positive examples $(y = 1)$

Negative examples $(y = 0)$

Photo OCR pipeline

- 1. Text detection

- 2. Character segmentation

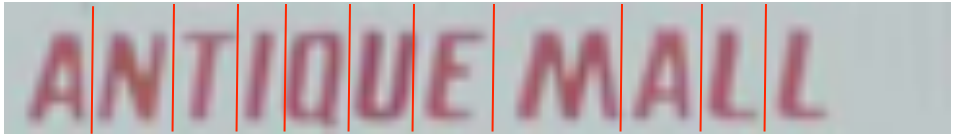

- 3. Character classification

just apply some supervised learning algorithm

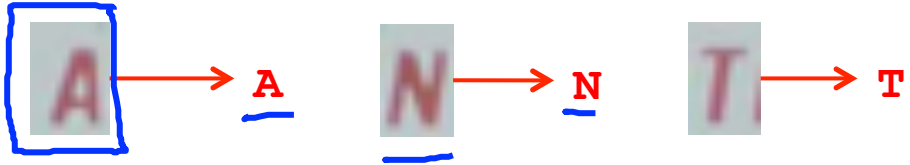

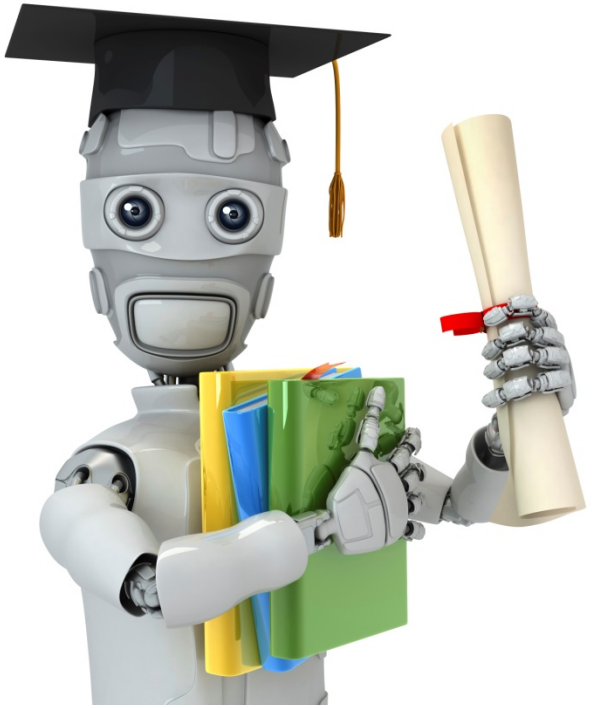

Machine Learning

Application example: Photo OCR

Getting lots of data: Artificial data synthesis

two main variations

Character recognition

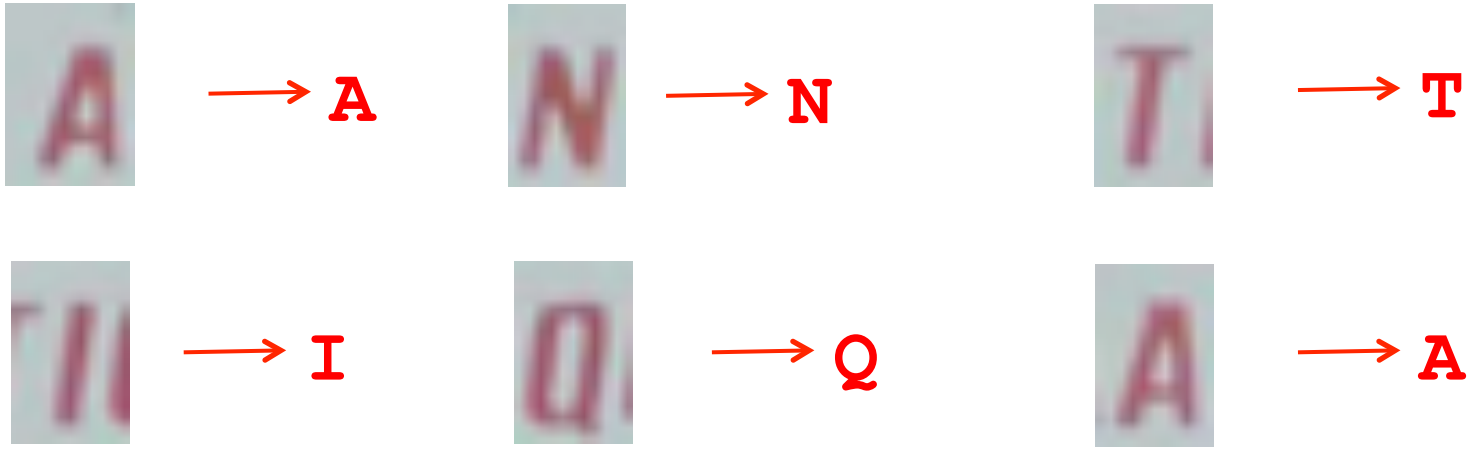

Artificial data synthesis for photo OCR

Real data

Abcdefg
Abcdefg
Abcdefg
Abcdefg
Abcdefg

[Adam Coates and Tao Wang]

Artificial data synthesis for photo OCR

Real data → Synthetic data

Creating new data from scratch

[Adam Coates and Tao Wang]

Synthesizing data by introducing distortions

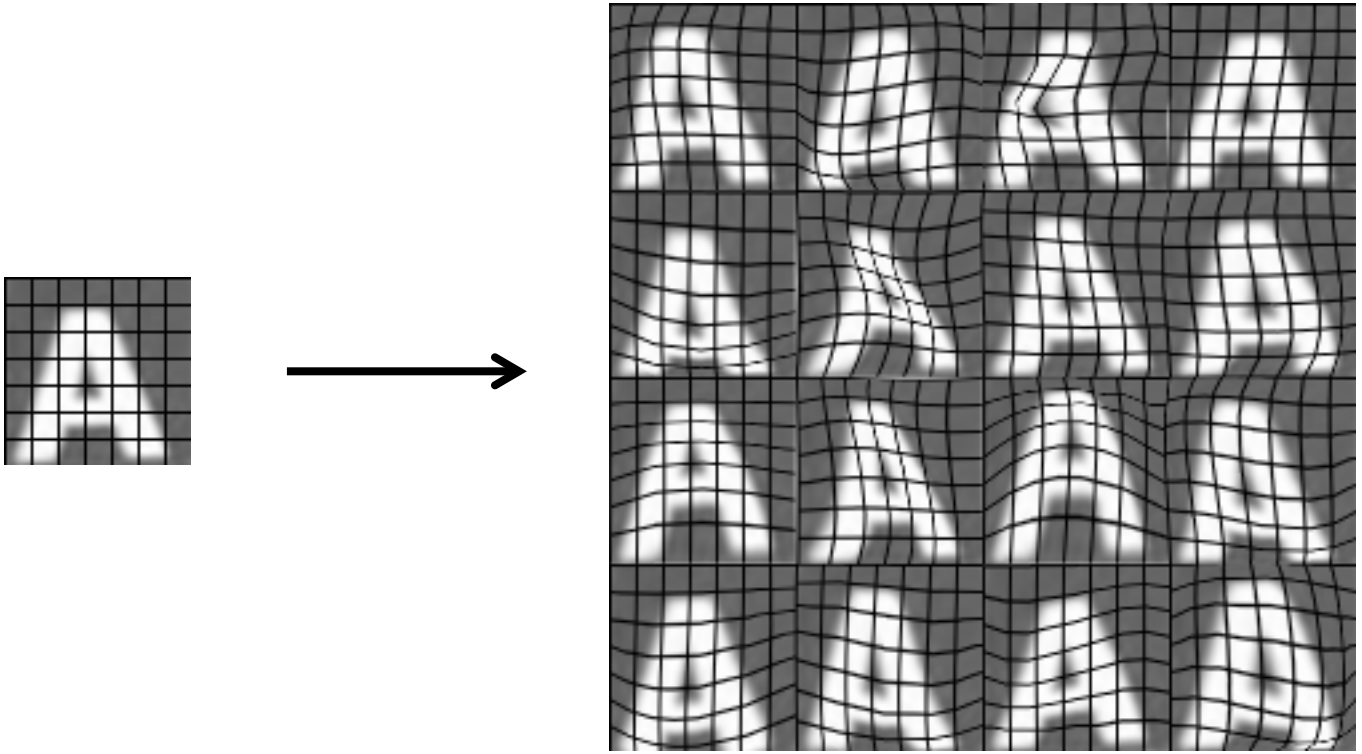

[Adam Coates and Tao Wang]

Synthesizing data by introducing distortions: Speech recognition

Original audio:

Audio on bad cellphone connection

Noisy background: Crowd

Noisy background: Machinery

[www.pdsounds.org]

Synthesizing data by introducing distortions

Distortion introduced should be representation of the type of noise/distortions in the test set.

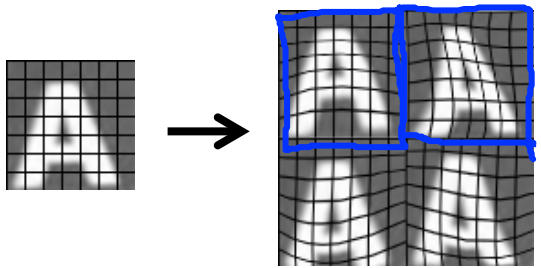

Audio:
Background noise,
bad cellphone connection

Usually does not help to add purely random/meaningless noise to your data.

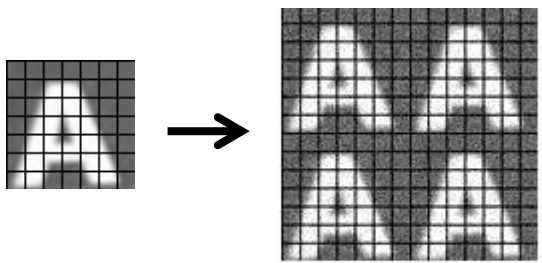

x_i = intensity (brightness) of pixel i

$x_i \leftarrow x_i +$ random noise

meaningless noise is less meaningful

[Adam Coates and Tao Wang]

Discussion on getting more data

1. Make sure you have a low bias classifier before expending the effort. (Plot learning curves). E.g. keep increasing the number of features/number of hidden units in neural network until you have a low bias classifier.
2. "How much work would it be to get 10x as much data as we currently have?"
   - Artificial data synthesis
   - Collect/label it yourself
   - "Crowd source" (E.g. Amazon Mechanical Turk)

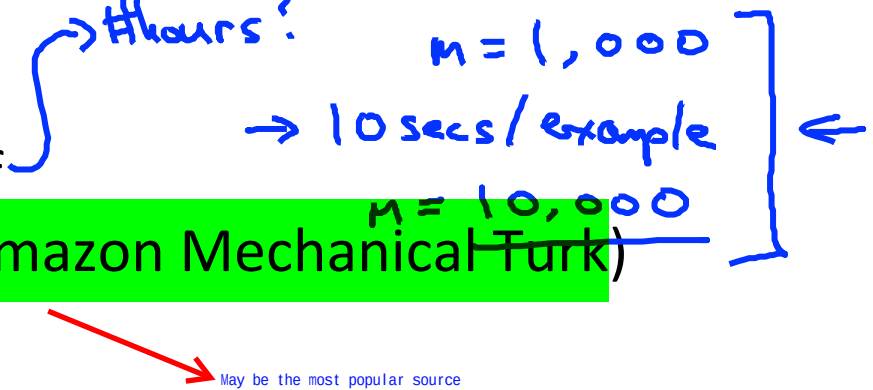

Discussion on getting more data

1. Make sure you have a low bias classifier before expending the effort. (Plot learning curves). E.g. keep increasing the number of features/number of hidden units in neural network until you have a low bias classifier.
2. "How much work would it be to get 10x as much data as we currently have?"
   - Artificial data synthesis
   - Collect/label it yourself
   - "Crowd source" (E.g. Amazon Mechanical Turk)

众包

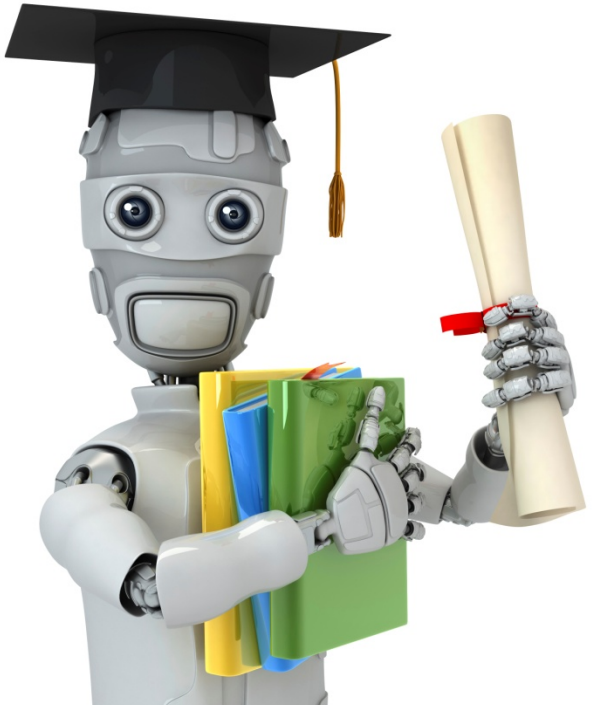

Application example: Photo OCR

Ceiling analysis: What part of the pipeline to work on next

Estimating the errors due to each component (ceiling analysis)

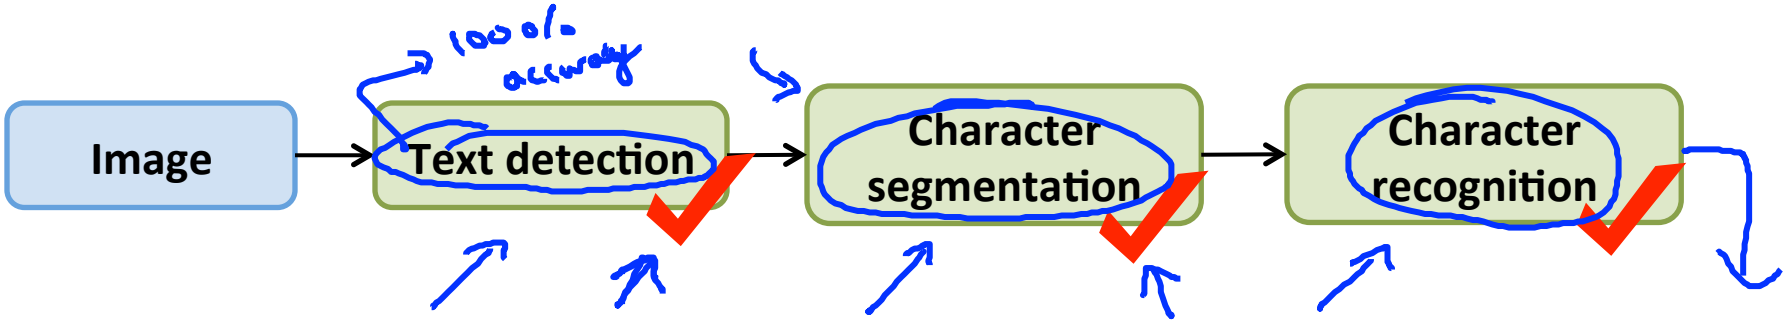

What part of the pipeline should you spend the most time trying to improve?

Component	Accuracy
Overall system	72%
Text detection (manually set that all text are correctly detected)	89%
Character segmentation (same idea)	90%
Character recognition (same)	100%

17%

1%

only 1% improvement (do not spend too much time here!)

10%

Another ceiling analysis example

Face recognition from images

(Artificial example)

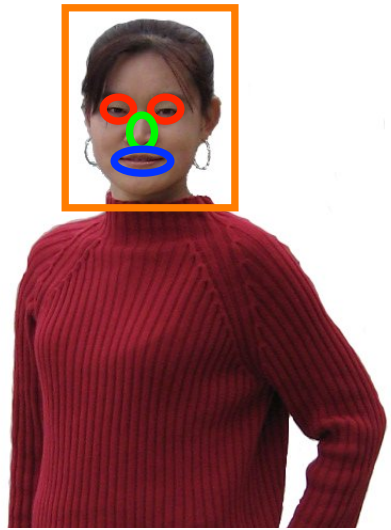

Another ceiling analysis example

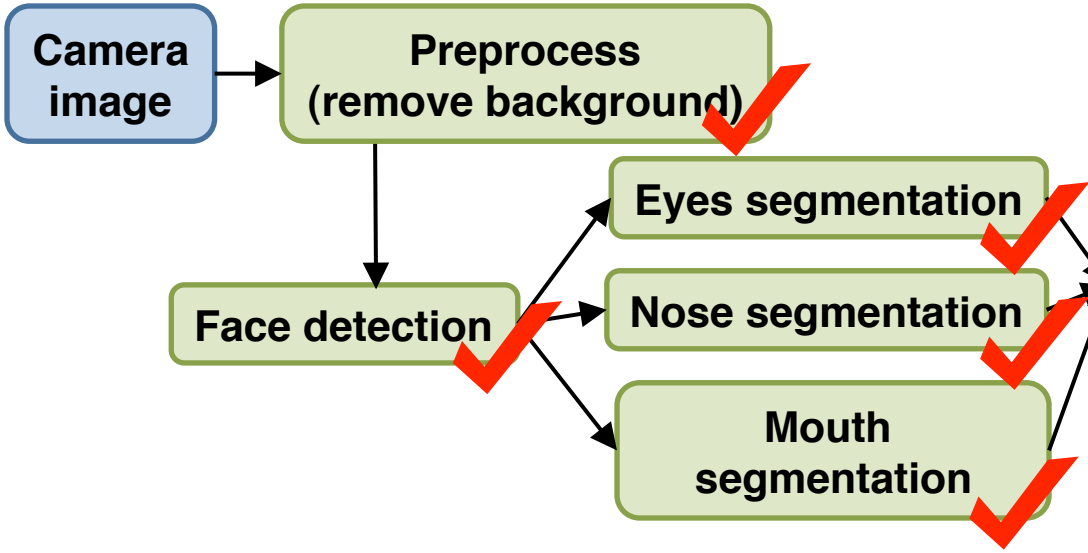

Component	Accuracy
Overall system	85%
Preprocess (remove background)	85.1%
Face detection	91%
Eyes segmentation	95%
Nose segmentation	96%
Mouth segmentation	97%
Logistic regression	100%

0.1%
5.9%
4%
1%
1%
3%